

Apartment Rental Prediction System

Dr. Ivan S. Zapreev

2020-01-07

Contents

Introduction	1
Dataset overview	1
Project goal	2
Execution plan	2
Data wrangling	2
Data cleaning	3
Initial steps	3
The main columns	3
Additional steps	7
Clean data summary	7
Splitting data	8
Data analysis	8
Modeling approach	8
Results	8
Conclusions	8
Future work	8
Appendix A: The complete list of data set columns	8
Appendix B: Data set column descriptions	9

Introduction

Dataset overview

As stated on the webpage of the ‘Apartment rental offers in Germany’ dataset, it contains 198,379 rental offers scraped from the Germany’s biggest real estate online platform β ImmobilienScout24.

The data set consists of a single CSV file: *immo_data.csv* which only contains offers for rental properties. The data features important rental property attributes, such as the living area size, the rent (both base rent as well as total rent), the location, type of energy, and etc. The **date** column present in the data set defines the time of scraping, which was done on three distinct dates: *2018-09-22*, *2019-05-10* and *2019-10-08*.

The complete list of data set columns is extensive¹ and thus in this study we will use the following subset:

¹Please consider reading “Appendix A” for the complete list of the data set columns.

```
## [1] "hasKitchen"      "heatingType"      "balcony"
## [4] "lift"            "garden"            "cellar"
## [7] "noParkSpaces"    "livingSpace"       "typeOfFlat"
## [10] "noRooms"         "floor"             "numberOfFloors"
## [13] "condition"       "newlyConst"        "interiorQual"
## [16] "yearConstructed" "energyEfficiencyClass" "regio1"
## [19] "regio2"          "regio3"            "baseRent"
## [22] "electricityBasePrice" "heatingCosts"      "serviceCharge"
## [25] "totalRent"       "date"
```

This sub-selection reduces the number of considered data set columns² from 48 to 26 and is motivated by the personal preferences of the report’s author and has no scientifically proven motivation. On the contrary, this column selection shall be seen as a part of problem statement. In other words, the task is to build an accurate³ rental price prediction model based on the predictors from this set of columns.

The additional data preparation steps will be described in the “*Data wrangling*” section of this document.

Project goal

Execution plan

Data wrangling

In this section we present cleaning, restructuring and enriching the raw data taken from the ‘Apartment rental offers in Germany’ dataset.

First, let us note that the number of data entries in the original data set is equal to 198332. This data is however not ready to be worked with as it contains multiple N/A values and other inconsistencies. For example, the next table summarizes the number of N/A values per column:

```
## # A tibble: 26 x 3
##   `Column name`      `N/A count` `N/A percent`
##   <chr>              <int>         <dbl>
## 1 electricityBasePrice 151158         76.2
## 2 energyEfficiencyClass 143315         72.3
## 3 heatingCosts         135154         68.2
## 4 noParkSpaces         130405         65.8
## 5 interiorQual         83001          41.8
## 6 numberOfFloors       71792          36.2
## 7 condition            50317          25.4
## 8 yearConstructed      42293          21.3
## 9 floor               37612          19.0
## 10 heatingType         32605          16.4
## 11 totalRent           29762          15.0
## 12 typeOfFlat          27571          13.9
## 13 serviceCharge        5110           2.58
## 14 hasKitchen           1              0
## 15 lift                 1              0
## 16 garden               1              0
## 17 cellar               1              0
## 18 livingSpace          1              0
## 19 noRooms               1              0
```

²Please consider reading “*Appendix B*” for the column descriptions.

³Please consider reading the “*Project goal*” section for an exact goal formulation.

## 20	baseRent	1	0
## 21	balcony	0	0
## 22	newlyConst	0	0
## 23	regio1	0	0
## 24	regio2	0	0
## 25	regio3	0	0
## 26	date	0	0

As one can see, about $\frac{1}{2}$ of the columns has 10–80% N/A^s, whereas the other half has (almost) no N/A^s.

The remainder of the section will be organized as follows. First we explain how we cleaned the data and solved some of its inconsistencies. Then we provide a summary of the cleaned data set. In the end we explain how we split the entire data set into the **validation** and **modeling** sub-sets⁴.

Data cleaning

The data cleaning will be explained in the next steps:

1. We begin with the **totalRent** column as this is the value that we want to predict;
2. We proceed with the columns with the marginal (< 1%) of N/A values;
3. We cover the remaining columns in the descending order of the number of N/A values.
4. We consider and solve some other data inconsistencies.

Initial steps

The **totalRent** column contains data that we want to predict. Therefore, the rows with **totalRent** == N/A are useless to us and shall be removed. Unfortunately, this will reduce the data set by 15.01%. There are also 13 columns with a marginal (0 to 1) number of N/A values. The latter can be seamlessly removed as even if all of these N/A^s appear in different rows, we will remove at most 13 entries which is just 0.0066% of data.

The main columns

Let us consider the columns one by one. Note that, some modifications we will do to the data to remove the N/A values may introduce bias. To for test that we would need a clean data set with no N/A values initially present and then to use such a data set for the trained model(s) validation. Due to the lack of time this will not be done in the case study.

Column: **electricityBasePrice** - 76.2% N/A values

We will set the electricity base price for the N/A values to zero. The motivation is that, since the number of N/A values is almost 80% and no other zero values are present:

```
x <- arog_data$selected_data %>% filter(!is.na(electricityBasePrice))
sum(x$electricityBasePrice == 0)
```

```
## [1] 0
```

it is likely that the N/A values were used to determine the fact that there is no electricity base price.

⁴The latter will also be split into the **training** and **testing** set for the sake of model cross-validation.

Column: energyEfficiencyClass - 72.3% N/A values

The energy efficiency factor levels are:

```
levels(aarg_data$selected_data$energyEfficiencyClass)
```

```
## [1] ""          "A"          "A_PLUS"     "B"
## [5] "C"         "D"          "E"          "F"
## [9] "G"         "H"          "NO_INFORMATION"
```

So we shall naturally set all the N/A and "" energy efficiency levels to "NO_INFORMATION".

Column: heatingCosts - 68.2% N/A values

We will set the heating costs for the N/A values to zero as there are already 1989 zero-valued heating cost entries. It is unlikely that there are non-heated accommodations in Germany so *we assume that the 0 values, the same as N/A's mean - "unknown"*.

Column: noParkSpaces - 65.8% N/A values

We will set the number of parking places for the N/A values to zero as there is already 2850 zero-valued entries. By this step we assume that, N/A is interpreted as *"not applicable"* or *"no are available"*.

Column: interiorQual - 38.8% N/A values

The interior quality factor levels are:

```
levels(aarg_data$selected_data$interiorQual)
```

```
## [1] ""          "luxury"     "normal"     "simple"
## [5] "sophisticated"
```

So we shall introduce a new level for the N/A and "" values, called "unknown".

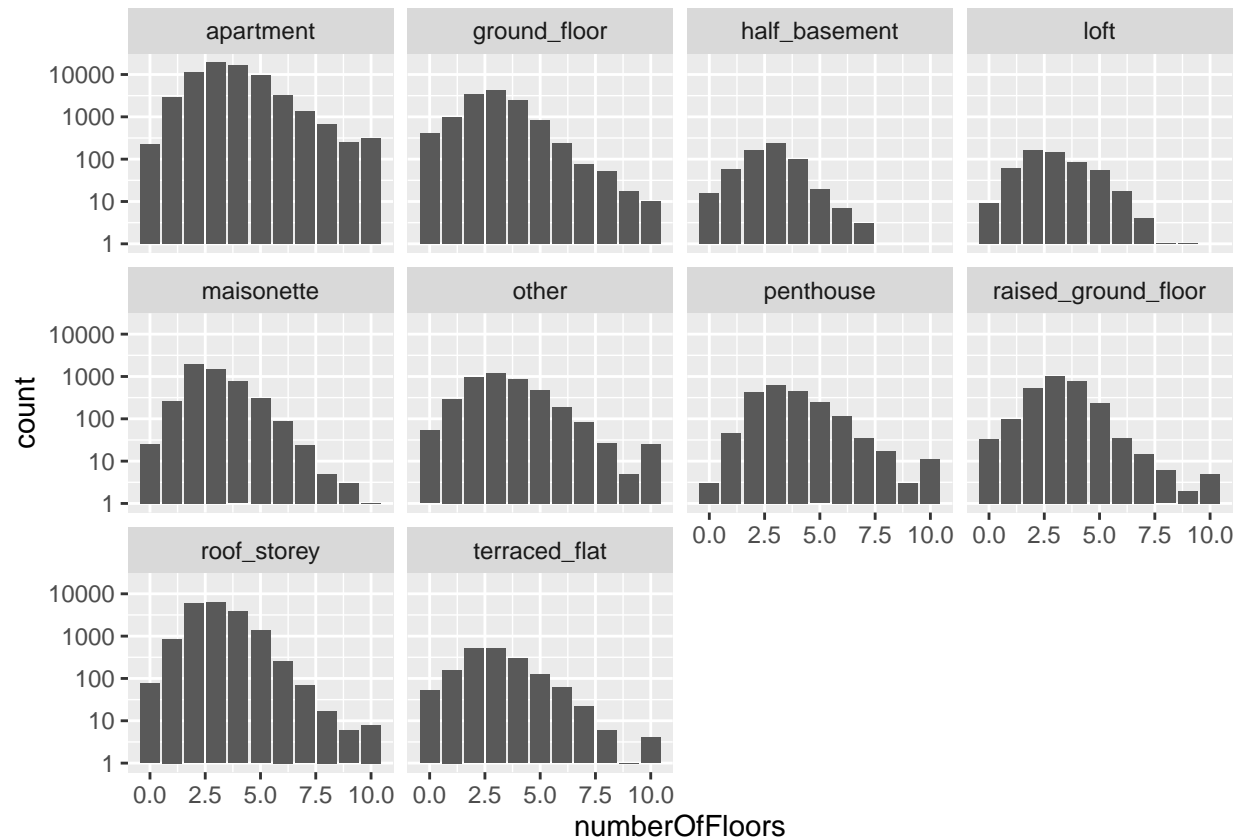
Column: numberOfFloors - 36.2% N/A values

Setting the N/A values for the floors shall be agreed with the apartment type, if we gather some number of floors statistics for each available apartment type we get the following:

Table 1: Type of flat vs. number of floors statistics

typeOfFlat	numberOfFloors				
	Count	Average	Standard error	Minimum	Maximum
apartment	66844	3.9	6.6	0	999
roof_storey	18450	3.1	6.7	0	800
ground_floor	12802	3	8.9	0	999
maisonette	4979	2.9	1.5	0	43
other	4284	3.6	5	0	301
raised_ground_floor	2779	3.4	7.3	0	370
penthouse	2011	3.7	2.2	0	33
terraced_flat	1760	3	1.5	0	14
half_basement	598	2.7	1.1	0	7
loft	542	3	1.5	0	15
""	0	NaN	NA	Inf	-Inf

From where we conclude that the data we have is very polluted. Clearly, one can not expect apartments with 99 floors and alike. See also on the large average (all +/- around 3 floors) and the huge standard error values. If we visualize the results (filtering out 1662 flats with more than 10 floors), we see that:



The data seems to be normally distributed (except for the **apartment** type) with the mean values within 2.5 - 4.0 range. This makes us believe that this data is too much biased and polluted. So we will not rely on this column in our analysis.

Column: condition - 25.4% N/A values

The condition factor levels are:

```
## [1] "" "first_time_use"
## [3] "first_time_use_after_refurbishment" "fully_renovated"
## [5] "mint_condition" "modernized"
## [7] "need_of_renovation" "negotiable"
## [9] "refurbished" "ripe_for_demolition"
## [11] "well_kept"
```

So we shall introduce a new level for the N/A and "" values, called "unknown".

Column: yearConstructed - 21.3% N/A values

There is no good default to replace the N/A values here. Yet, it is a significant amount of data which we do not want to exclude. Therefore drop this column from the analysis and just use the **newlyConst** flag column.

Column: floor - 19.0% N/A values

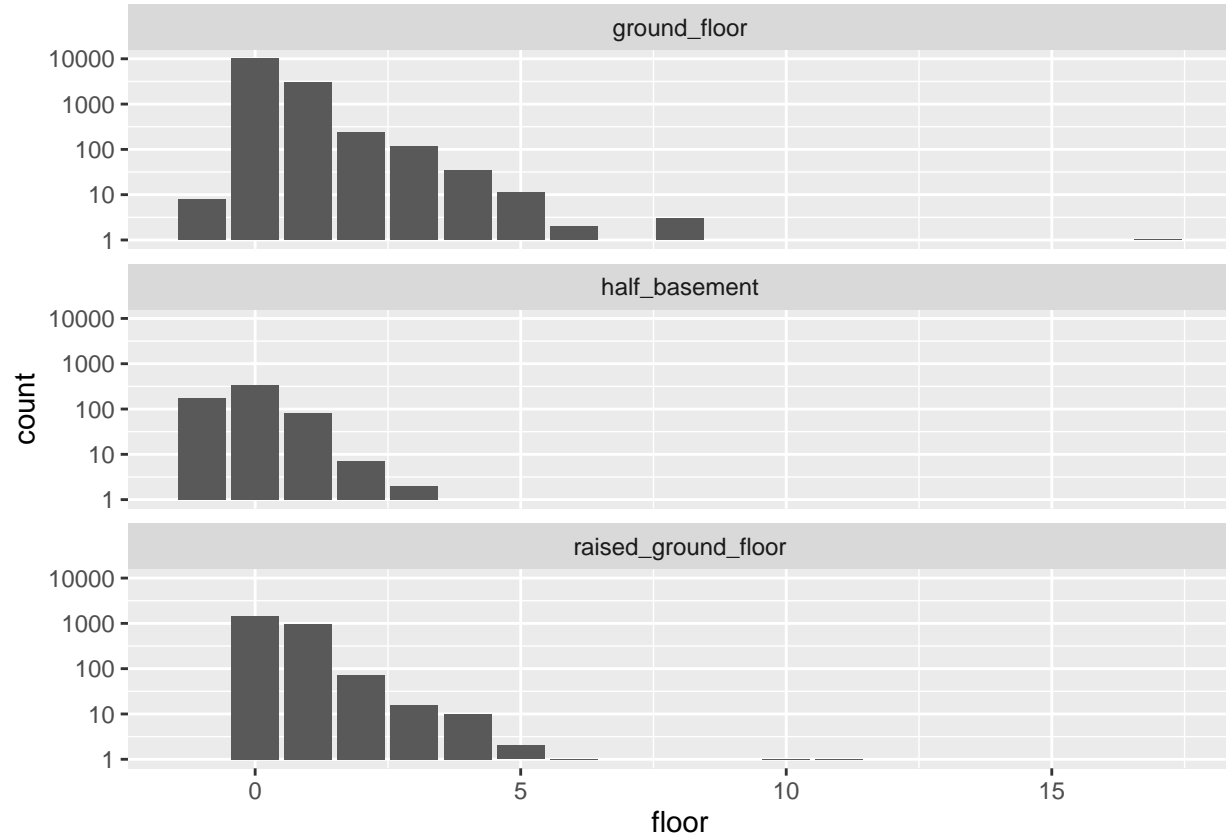
We could assign some **floor** values based on the flat types:

```
## [1] "" "apartment" "ground_floor"
## [4] "half_basement" "loft" "maisonette"
## [7] "other" "penthouse" "raised_ground_floor"
## [10] "roof_storey" "terraced_flat"
```

For example, we could consider assigning:

- `half_basement` – the average floor for the half-basement
- `ground_floor` – floor = 0
- `raised_ground_floor` – the average floor for the raised ground floor

but, let us look at the floor values (filtering out 10 flats higher than at the 100'th floor), for these flat types:



From the data above we see that we shall not only correct the N/A values but set all of the floor values for the considered flat types as follows:

- `half_basement` – floor = -1
- `ground_floor` – floor = 0
- `raised_ground_floor` – floor = 0

If we do that then there will still be 20050 (10.1% of data) N/A floor values for the flat types for which we can not give any exact value. So we will just assign those to the mean floor value in the category.

Column: `heatingType` - 16.4% N/A values

The heating type factor levels are:

```
## [1] "central_heating" "combined_heat_and_power_plant"
## [3] "district_heating" "electric_heating"
## [5] "floor_heating" "gas_heating"
## [7] "heat_pump" "night_storage_heater"
```

```
## [9] "oil_heating"          "self_contained_central_heating"
## [11] "solar_heating"        "stove_heating"
## [13] "wood_pellet_heating"
```

So we shall introduce a new level for the N/A, "", and "H" values, called "unknown".

Column: typeOfFlat - 13.9% N/A values

The type of flat factor levels are:

```
## [1] ""          "apartment"    "ground_floor"
## [4] "half_basement" "loft"         "maisonette"
## [7] "other"        "penthouse"    "raised_ground_floor"
## [10] "roof_storey"  "terraced_flat"
```

So we shall introduce a new level for the N/A and "" values, called "unknown". Note that, we do not use the pre-defined level "other" here as we interpret it as known flat type which is just not on the list of available choices.

Column serviceCharge - 2.58% N/A values

We will set the service charges for the N/A values to zero. The motivation is that, there are:

```
x <- arog_data$selected_data %>% filter(!is.na(serviceCharge))
sum(x$serviceCharge == 0)
```

```
## [1] 2496
```

zero values present, so we interpret the N/A values as defining the fact of no additional service charges.

Additional steps

In addition to the data alternations done above we have also done the following:

- Re-setting the number of floors:
 - half_basement - floor = -1
 - ground_floor - floor = 0
 - raised_ground_floor - floor = 0
- Filter out flats:
 - With floor > 100
 - Other than "half_basement", "other", and "unknown"; but with floor < 0

Clean data summary

Let us now summarize the resulting clean data:

```
## # A tibble: 24 x 3
##   `Column name`      `N/A count` `N/A percent`
##   <chr>             <int>         <dbl>
## 1 hasKitchen         0             0
## 2 heatingType        0             0
## 3 balcony           0             0
## 4 lift              0             0
## 5 garden            0             0
## 6 cellar            0             0
## 7 noParkSpaces      0             0
```

## 8 livingSpace	0	0
## 9 typeOfFlat	0	0
## 10 noRooms	0	0
## 11 floor	0	0
## 12 condition	0	0
## 13 newlyConst	0	0
## 14 interiorQual	0	0
## 15 energyEfficiencyClass	0	0
## 16 regio1	0	0
## 17 regio2	0	0
## 18 regio3	0	0
## 19 baseRent	0	0
## 20 electricityBasePrice	0	0
## 21 heatingCosts	0	0
## 22 serviceCharge	0	0
## 23 totalRent	0	0
## 24 date	0	0

As one can notice, the dat set size has been reduced from 198332 to 164637 . The major reason for that is excluding the rows with the N/A values of the `totalRent` column. Let us recall that the number of such raws was 15.01% of the data set, e.g. 29770 rows. It now remains to notice that $198332 - 29770 = 168562 \approx 164637$. The remaining 2% delta is explained by cleaning the `floor/typeOfFlat` columns and etc.

Splitting data

Data analysis

Modeling approach

Results

Conclusions

Future work

Check for introducing any bias by data wrangling.

Appendix A: The complete list of data set columns

Hereby we present the list of columns from the original data set:

## [1] "regio1"	"serviceCharge"
## [3] "heatingType"	"telekomTvOffer"
## [5] "telekomHybridUploadSpeed"	"newlyConst"
## [7] "balcony"	"electricityBasePrice"
## [9] "picturecount"	"pricetrend"
## [11] "telekomUploadSpeed"	"totalRent"
## [13] "yearConstructed"	"electricityKwhPrice"
## [15] "scoutId"	"noParkSpaces"

## [17] "firingTypes"	"hasKitchen"
## [19] "geo_bln"	"cellar"
## [21] "yearConstructedRange"	"baseRent"
## [23] "houseNumber"	"livingSpace"
## [25] "geo_krs"	"condition"
## [27] "interiorQual"	"petsAllowed"
## [29] "streetPlain"	"lift"
## [31] "baseRentRange"	"typeOfFlat"
## [33] "geo_plz"	"noRooms"
## [35] "thermalChar"	"floor"
## [37] "numberOfFloors"	"noRoomsRange"
## [39] "garden"	"livingSpaceRange"
## [41] "regio2"	"regio3"
## [43] "description"	"facilities"
## [45] "heatingCosts"	"energyEfficiencyClass"
## [47] "lastRefurbish"	"date"

Appendix B: Data set column descriptions

Here is the list of the initially considered data set columns with the descriptions thereof:

1. `hasKitchen` – has a kitchen
2. `balcony` – does the object have a balcony
3. `cellar` – has a cellar
4. `lift` – is elevator available
5. `floor` – which floor is the flat on
6. `garden` – has a garden
7. `noParkSpaces` – number of parking spaces
8. `livingSpace` – living space in sqm
9. `condition` – condition of the flat
10. `interiorQual` – interior quality
11. `regio1` – Bundesland
12. `regio2` – District or Kreis, same as geo krs
13. `regio3` – City/town
14. `noRooms` – number of rooms
15. `numberOfFloors` – number of floors in the building
16. `typeOfFlat` – type of flat
17. `yearConstructed` – construction year
18. `newlyConst` – is the building newly constructed
19. `heatingType` – Type of heating
20. `energyEfficiencyClass` – energy efficiency class
21. `heatingCosts` – monthly heating costs in €

- 22. `serviceCharge` – auxiliary costs such as electricity or Internet in €
- 23. `electricityBasePrice` – monthly base price for electricity in €
- 24. `baseRent` – base rent without electricity and heating
- 25. `totalRent` – total rent (usually a sum of base rent, service charge and heating cost)
- 26. `date` – time of scraping